

BoothPilot

Your Booth's AI Demo Concierge

An interactive conversational-AI booth that gives every visitor a personalized live product demo — and turns it into pipeline.

AI Growth Hackathon · Orange Slice × Y Combinator · Project Brief

Live booth conversation

Visitor: “I run growth at a fintech — our problem is churn.”

BoothPilot: “Let me show you the churn view on your data shape.” *[opens the live product]*

Visitor: “Can it flag at-risk accounts weekly?”

BoothPilot: “Yes — here. Want me to send this to your team?”

91

Pursuit Score

Persona: Growth lead
Intent: High · Fit: Strong
Next: Send pilot follow-up

1 Executive Summary

Core product. BoothPilot is an interactive, conversational-AI booth assistant. When a visitor walks up, it greets them, holds a real two-way conversation, and **drives the company's actual product live** — tailoring the demo to that person's exact use case. In the background it scores how serious the visitor is, captures the full context, and sends a personalized follow-up. Every visitor gets a senior-rep-quality demo; the company gets a scored pipeline.

Simple explanation. At a conference, a booth has one or two people and a line of curious visitors. Most get a generic pitch or a quick brush-off, and the good ones are forgotten by Monday. BoothPilot gives *each* visitor their own personalized, hands-on product demo at the same time — and remembers all of them.



simultaneous personalized demos

0

visitors waiting in line

100%

conversations captured & scored

Sales Cyborgs

Zero to One

Reading Minds

Revenue on Autopilot

Algorithm Hacking

2 The Problem: Booths Don't Scale

A booth is one of the most expensive things a startup does — the floor space, the travel, the team's time. The job is simple: turn curious foot traffic into pipeline. Today that breaks in three places.

- **One rep, many visitors.** While one person gets a demo, three others wait, half-listen, and wander off. Attention is the scarcest thing on the floor and it's rationed badly.
- **One pitch for everyone.** A founder, a buyer, an engineer, and a tire-kicker all hear the same scripted spiel. The demo almost never matches the visitor's actual use case — which is the only thing that makes them lean in.
- **Follow-up rot.** Notes are written later, if at all. The follow-up is late, generic, or never sent. Good leads blur together with low-value contacts.

The real loss: every visitor needed a *different* demo, and got the same one — or none. The company met people and never showed each of them the one thing that would have made them care.

3 Why Not Just a Human? Why Not ChatGPT?

This is the first question a sharp judge asks. BoothPilot has a clean answer to each.

The objection	Why BoothPilot wins
“Just put a person there.”	A human gives <i>one</i> demo at a time and forgets the details. BoothPilot gives <i>every</i> visitor a tailored, hands-on, live-product walkthrough simultaneously — and remembers all of them perfectly. Humans don’t scale; this does.
“I’ll just ask ChatGPT on my phone.”	ChatGPT can’t operate the company’s real platform, can’t show <i>their</i> features on <i>your</i> use case, and — critically — doesn’t drop you into the company’s pipeline. BoothPilot turns the conversation into a captured, scored, followed-up lead.
“Isn’t this a pre-recorded demo?”	No. It’s agentic: it listens, asks, and then <i>drives the live product</i> to whatever the visitor cares about. The path is different for every person.

The wedge: personalized, live product demos at scale, with pipeline capture built in. No incumbent badge-scanner or chatbot does the live, agentic, product-driving demo — that is the defensible core.


4 The Solution

BoothPilot is a physical booth station running a conversational AI concierge with **agentic access to the company’s product**. It is deliberately two-sided.



The visitor never feels “scored.” They feel *helped*. The company gets the intelligence. Keeping these two surfaces separate is a core design rule — buyers want to make their own decision, not be told what to buy.

5 Product Flow



1. A visitor approaches; the booth proactively greets them (anonymous presence detection — see §12).
2. They talk naturally: who they are, what they do, what brought them over.
3. BoothPilot extracts the real use case and any objections through conversation, not a form.

4. It **opens the live product** and walks them through *their* scenario — the exact feature, on their data shape.
5. It offers to send the personalized walkthrough to the visitor and their colleagues, capturing emails as a natural part of the conversation.
6. In the background: persona, intent, fit, and a Pursuit Score are computed; a tailored follow-up email is drafted and sent.

6 The Personalized Demo Engine

This is the hero capability. BoothPilot is granted scoped, agentic access to the company’s product through tool-calling, and it operates that product live in response to the conversation.

- **Use-case-driven navigation.** “Can it flag at-risk accounts weekly?” → the agent opens the churn view, configures a weekly flag, and shows the result.
- **Tailored, not templated.** The demo path is generated per visitor from what they said — a founder, a buyer, and an engineer get three different tours of the same product.
- **Safe by design.** The agent drives a sandboxed instance with a curated set of tools and demo data, so it is impressive and reliable, never destructive.
- **Answers in context.** If the visitor asks about a feature, BoothPilot shows it rather than describing it — the single most persuasive thing a demo can do.

7 Real Example

Scenario. A company is exhibiting an analytics platform. **Maya Rodriguez**, Head of Growth at a fintech, walks up.

What happens at the booth

Conversation	“I run growth at a fintech; churn is killing us. We can’t see which accounts are slipping until they cancel.”
BoothPilot action	Opens the live product, builds a churn-risk board on a fintech-shaped sample, sets a weekly at-risk alert, and walks Maya through reading it.
Visitor takeaway	A personalized walkthrough link + the option to forward it to her two teammates who own retention.

What the company sees (invisible to Maya)

Pursuit Score	91 / 100 — strong ICP fit, high-intent language (“killing us”), 4-minute engaged demo, decision-maker.
Pain point	Reactive churn; no early-warning signal.
Best angle	Weekly at-risk alerting; offer a 2-week pilot on their data.
Next action	Pilot follow-up sent automatically; flag for AE intro.
Follow-up	Drafted and sent, referencing the exact churn board she saw.

Result: the company no longer has “a badge scan from Maya.” It has a demoed, scored, pilot-ready opportunity — and two of her teammates already invited in.

8 What the Visitor Gets

The visitor should leave feeling they got something genuinely useful — not that they filled out a lead form.

- A **personalized walkthrough** of the product on their own use case, not a generic tour.
- A **recap link** (and optional printed receipt) capturing exactly what they saw, which they can forward to their team.
- A reason to bring colleagues in: “Unlock your team’s version — we’ll tailor it to each of them.”


9 What the Company Gets

A live dashboard with a rich lead card for every visitor — far beyond name and email.

- **Pursuit Score** with a transparent, multi-factor reason (see §11).
- **Full conversation context:** their words, their use case, their objections.
- **What they were shown**, so the follow-up references the real moment.
- **Auto-sent, personalized follow-up** and a recommended next action.
- **Product-gap analytics:** “6 visitors today asked about SSO and we had no answer” — a growth and roadmap signal a human booth never captures.

10 The Growth Loop

A booth that only captures the people who walk up is a funnel bottom. BoothPilot turns the booth into a funnel *top*.



The forwarded recap is genuinely useful and makes the sender look good to their team — that is the share incentive. Referred colleagues did not have to be on the show floor, so reach escapes the booth.

The one number to repeat on stage: reach created per person who walks up. With a modest share rate, every booth visit pulls in additional qualified people the company would never have met — the same floor space, multiplied. We present the loop math honestly, including the share-rate assumptions, because rigor is the point.

11 The Pursuit Score (For the Company, Not the Visitor)

A teammate’s fair critique: buyers do not want an AI telling them what to buy. So the score is never shown to the visitor and the agent never plays gatekeeper. It is an *internal* signal, built from multiple transparent factors:

Factor	Signal
ICP fit	Role, company, industry — enriched and cross-checked against what they typed.
Intent language	Urgency and buying cues in their own words (“evaluating now,” “killing us”).
Engagement	Conversation depth and dwell time — behavioral, not self-reported.
Authority	Decision-maker vs. influencer vs. curious.
Demo depth	Which features they pushed on and for how long.

Every score ships with a human-readable reason, so the booth team can trust it and act on it.


12 Behavioral Signal Layer

An optional camera adds a behavioral dimension that forms cannot: it detects when someone is approaching (so the booth greets proactively) and measures engagement and dwell as an intent signal.

Privacy by design. Anonymous presence and engagement only — *detection, not recognition*. Opt-in, processed on-device, nothing stored, no identity. We pitch “the booth notices interest,” never “the booth recognizes you.” If in doubt, this layer is cut without affecting the core.


13 Hardware Concept

The physical station makes the experience real and memorable; a web form is easy to ignore.




14 Agentic AI Pipeline

The AI is not one prompt. It is a visible, multi-step agent — and the steps are shown on the dashboard so judges can see the depth.



- **Converse:** OpenAI Realtime voice for low-latency, natural two-way conversation.
- **Drive product:** function/tool-calling against a curated set of product actions on a sandboxed instance.
- **Score with reasoning:** structured outputs produce a Pursuit Score plus the factors behind it.
- **Act, don't draft:** the follow-up is actually *sent*, and a day-2 sequence is scheduled.
- **Flag gaps:** recurring unmet asks are aggregated into a roadmap/growth signal.

15 Technical Architecture



Frontend	Next.js — visitor view, live product frame, and the company dashboard.
Backend	Convex — realtime lead store, mutations/queries, and actions that orchestrate the agent (sponsor stack).
AI	OpenAI Realtime (voice) + agents with tool-calling and structured outputs (sponsor stack).
Demo target	A sandboxed sample product the agent drives safely on stage.
Email	Resend for actually sending the personalized follow-up.
Built with	Cursor (sponsor stack). Open-sourced on GitHub per the rules.

16 Hackathon Track Fit

Track	Why BoothPilot fits	Fit
Sales Cyborgs	An AI sales engineer that demos, qualifies, and writes follow-ups at scale.	High
Zero to One	Personalized, hands-on product onboarding for every prospect.	High
Reading Minds	Detects intent, fit, and engagement from conversation and behavior.	High
Revenue on Autopilot	Turns booth traffic into auto-sequenced pipeline.	Medium
Algorithm Hacking	Forwardable personalized demos create reach beyond the floor.	Medium

The strongest story is not breadth. It is that BoothPilot *dominates* Sales Cyborgs and Zero to One: it is the one project that gives every visitor a real, personalized product experience and converts it into pipeline.

17 MVP Scope for 24 Hours

Do not overbuild. The core demo must be reliable and must land the “personalized live demo” wow.

Build first	Voice conversation, agent driving the sandboxed product live, intent + Pursuit Score, dashboard, auto-sent follow-up.
Then	Forwardable recap (growth loop), printed receipt, proactive greet.
Optional polish	Camera engagement signal, branded enclosure, product-gap analytics view, CRM export.
Cut if behind	Anything that is not the live personalized demo. That moment is the whole pitch.

Required MVP: (1) voice conversation, (2) agentic live product demo, (3) intent + Pursuit Score with reason, (4) live dashboard, (5) auto-sent personalized follow-up, (6) phone fallback.

18 24-Hour Build Plan

Golden rule: build the spine first. By hour 4 have *voice in* → *agent opens one product view* → *lead card on dashboard* working end to end, even if stubbed. Everything after enriches a working demo.

h0–4 (Sat eve)	Repo public; Next.js + Convex wired; OpenAI Realtime voice loop; sandbox product target chosen; stubbed score on dashboard. <i>Checkpoint: end-to-end stub works.</i>
h4–10 (overnight)	Real agent tool-calling drives the product; intent + Pursuit Score with reasons; follow-up actually sends via Resend. <i>Checkpoint: full pipeline live.</i>
h10–16 (early AM)	Forwardable recap + growth-loop attribution; proactive greet; reliability + latency tuning; pre-seeded golden sessions.
h16–21 (Sun)	Stage polish; rehearse the demo 3×; record a backup video; edge cases.
h21–judging	Code freeze, booth setup, pre-warm everything.

Team split: Frontend (visitor + dashboard) · Backend/AI (Convex + agent pipeline) · Product/hardware (sandbox target + Pi/voice) · Demo & polish (script, reliability, pitch).

19 The Winning Demo (2–3 min)

De-risk first: pre-seed sessions, pre-warm OpenAI + Convex, a planted “colleague” for the loop, a recorded backup, and tight timeouts with cached fallbacks so there is never dead air.

1. **Hook line:** “Every booth here can pitch one person at a time. Watch this one give a personalized demo to everyone at once.”
2. A judge walks up and *talks* to the booth: their role, their problem.
3. **Wow #1:** the booth opens the real product and demos *their* use case, live, by voice.
4. **Wow #2:** “Send this to my team” → a colleague’s session appears on the dashboard, auto-scored, follow-up sent. Reach climbs on screen.
5. Show the company side: Pursuit Score 91 with its reasons, the drafted follow-up, the product-gap signal.
6. **Closing line:** “Other booths hand you a spreadsheet on Monday. This one hands every visitor their own demo — and hands you a pipeline. It’s open-source and runs on a \$40 Pi.”

20 Pitch & Preempting Critiques

One-liner. BoothPilot is a conversational-AI booth that gives every visitor a personalized live demo of your product — and turns each conversation into a scored, followed-up lead.

Longer pitch. Companies spend thousands on a booth and a line of curious visitors, but one rep can only demo one person at a time, the pitch is generic, and follow-up rots. BoothPilot is an AI concierge with agentic access to your product: it converses, demos your live platform on each visitor’s use case, scores intent in the background, and sends a personalized follow-up — while turning each visit into referred reach beyond the floor.

Preempt the two killer critiques head-on:

- “*It’s just a chatbot.*” “A chatbot talks. This *drives our live product* on your use case and sends a follow-up referencing exactly what you saw.”

- “*It’s capture, not growth.*” “Capture improves traffic you have. We give every visitor a demo a human couldn’t, and the recap they forward pulls in people who were never on the floor.”

21 What We Are Not Building

BoothPilot is not a generic chatbot, not a badge-scanning form, and *not* an AI that decides for the buyer. It never tells a visitor what they should purchase. It helps them explore, and it helps the company know who to pursue.

22 Success Criteria

By judging time the product should show: (1) a real voice conversation; (2) the agent driving the live product on a visitor’s use case; (3) an intent-based Pursuit Score with reasons; (4) an auto-sent personalized follow-up; (5) the growth loop adding a referred session on screen; (6) a reliable phone fallback; (7) a crisp pitch.

Final framing. BoothPilot is what happens when a sales engineer, a product demo, and a CRM merge into one AI concierge that works every visitor at once. It answers the question every booth team has on Monday: *who did we actually move, and what should we say next?*

End of project brief